

Python code to implement bagging using Random Forest and boosting using AdaBoost and Gradient Boosting.

Explanation

Loading Dataset: `load_iris()`: Loads the Iris dataset, a popular dataset for classification tasks. X and y: Features and target variables.

Splitting Dataset: `train_test_split()`: Splits the dataset into training and testing sets with a specified test size (e.g., 20%).

Model Initialization: `RandomForestClassifier()`: Initializes the Random Forest model with a specified number of estimators.

`DecisionTreeClassifier(max_depth=1)`: Initializes a simple decision tree (decision stump) for AdaBoost. `AdaBoostClassifier()`: Initializes the AdaBoost model with the decision stump as the base estimator and a specified number of estimators. `GradientBoostingClassifier()`: Initializes the Gradient Boosting model with specified parameters like the number of estimators, learning rate, and max depth.

Training the Model: `fit()`: Trains the model on the training data.

Making Predictions: `predict()`: Makes predictions on the testing data.

Evaluating the Model: `accuracy_score()`: Calculates the accuracy of the model. `classification_report()`: Provides a detailed classification report with precision, recall, and F1-score. `confusion_matrix()`: Generates the confusion matrix to visualize the performance.

Bagging with Random Forest

Step-by-Step Implementation Import Libraries: Load necessary libraries. Load Dataset: Use a sample dataset from scikit-learn. Split Dataset: Split the data into training and testing sets. Initialize and Train Model: Set up the Random Forest model and train it. Predict and Evaluate: Make predictions and evaluate the model.

```
In [4]: # Import necessary Libraries
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.datasets import load_iris
import pandas as pd

# Load dataset
data = load_iris()
X = data.data
y = data.target
```

```
# Display dataset information
df = pd.DataFrame(X, columns=data.feature_names)
df['target'] = y
print("First 5 rows of the dataset:")
print(df.head())

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the RandomForestClassifier
model = RandomForestClassifier(n_estimators=100, random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Random Forest Accuracy: {accuracy:.2f}')

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

First 5 rows of the dataset:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

Random Forest Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Boosting with AdaBoost

Step-by-Step Implementation Import Libraries: Load necessary libraries. Load Dataset: Use a sample dataset from scikit-learn. Split

Dataset: Split the data into training and testing sets. Initialize and Train Model: Set up the AdaBoost model with a weak learner

(decision stump) and train it. Predict and Evaluate: Make predictions and evaluate the model.

```
In [5]: # Import necessary Libraries
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.datasets import load_iris
```

```
import pandas as pd

# Load dataset
data = load_iris()
X = data.data
y = data.target

# Display dataset information
df = pd.DataFrame(X, columns=data.feature_names)
df['target'] = y
print("First 5 rows of the dataset:")
print(df.head())

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the base estimator (decision stump)
base_estimator = DecisionTreeClassifier(max_depth=1)

# Initialize the AdaBoostClassifier
model = AdaBoostClassifier(base_estimator=base_estimator, n_estimators=100, random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'AdaBoost Accuracy: {accuracy:.2f}')

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

First 5 rows of the dataset:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

F:\Users\chaitra\anaconda3\Lib\site-packages\sklearn\ensemble_base.py:156: FutureWarning: `base_estimator` was renamed to `estimator` in version 1.2 and will be removed in 1.4.

warnings.warn(

AdaBoost Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

Boosting with Gradient Boosting

Step-by-Step Implementation Import Libraries: Load necessary libraries. Load Dataset: Use a sample dataset from scikit-learn. Split Dataset: Split the data into training and testing sets. Initialize and Train Model: Set up the Gradient Boosting model and train it. Predict and Evaluate: Make predictions and evaluate the model.

```
In [6]: # Import necessary Libraries
        from sklearn.ensemble import GradientBoostingClassifier
```

```
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.datasets import load_iris
import pandas as pd

# Load dataset
data = load_iris()
X = data.data
y = data.target

# Display dataset information
df = pd.DataFrame(X, columns=data.feature_names)
df['target'] = y
print("First 5 rows of the dataset:")
print(df.head())

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize the GradientBoostingClassifier
model = GradientBoostingClassifier(n_estimators=100, learning_rate=0.1, max_depth=3, random_state=42)

# Train the model
model.fit(X_train, y_train)

# Make predictions
y_pred = model.predict(X_test)

# Evaluate the model
accuracy = accuracy_score(y_test, y_pred)
print(f'Gradient Boosting Accuracy: {accuracy:.2f}')

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

First 5 rows of the dataset:

	sepal length (cm)	sepal width (cm)	petal length (cm)	petal width (cm)	\
0	5.1	3.5	1.4	0.2	
1	4.9	3.0	1.4	0.2	
2	4.7	3.2	1.3	0.2	
3	4.6	3.1	1.5	0.2	
4	5.0	3.6	1.4	0.2	

	target
0	0
1	0
2	0
3	0
4	0

Gradient Boosting Accuracy: 1.00

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

Confusion Matrix:

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

In []: